

Name : Vaibhav Gupta
Div : D15C
Batch : C
Roll No : 61

MLDL Experiment 01

Aim: Implement Linear and Logistic Regression on real-world datasets.

Linear Regression:

1. Dataset Source:

- **Dataset Name:** House Price Prediction Dataset
- **Source:**
<https://www.kaggle.com/datasets/suvidyasonawane/house-price-prediction-dataset>
- **Repository Owner:** Suvidya Sonawane

2. Dataset Description:

This dataset contains 300 synthetic housing records created for practicing regression and machine learning workflows.

The data simulates realistic relationships between housing features and sale prices, making it ideal for exploratory data analysis (EDA), feature engineering, and regression modeling.

- **Size:** 300 samples (rows) × multiple numerical attributes (columns)

- Target Variable:
 - SalePrice (Continuous) – Final selling price of the house
- Key Features:
 1. OverallQual – Overall material and finish quality of the house
 2. GrLivArea – Above ground living area (square feet)
 3. GarageCars – Number of cars that can fit in the garage
 4. TotalBsmtSF – Total basement area (square feet)
 5. YearBuilt – Year the house was constructed
 6. TotRmsAbvGrd – Total rooms above ground
 7. LotArea – Size of the lot in square feet

The dataset is clean, beginner-friendly, and requires no preprocessing, allowing direct application of regression algorithms.

3. Mathematical Formulation of the Algorithm:

Multiple Linear Regression models the relationship between several independent variables and a continuous dependent variable.

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

Where:

- \hat{y} = Predicted house price
- β_0 = Intercept
- $\beta_1 \dots \beta_n$ = Coefficients
- $x_1 \dots x_n$ = Input features

Cost Function (Mean Squared Error):

$$J(\beta) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

4. Algorithm Limitations:

1. **Linearity Assumption:** Assumes linear relationship between features and price.
2. **Sensitive to Outliers:** Extremely expensive houses may distort predictions.
3. **Multicollinearity:** Highly correlated features can reduce model stability.

5. Methodology / Workflow:

1. Load CSV dataset

2. Select input features and target (SalePrice)
3. Split dataset into training and testing sets (80:20)
4. Train Multiple Linear Regression model
5. Predict house prices
6. Evaluate using MSE and R² score

6. Performance Analysis:

- **Metric 1:** Mean Squared Error (MSE)
 - Measures average squared prediction error
- **Metric 2:** R-Squared (R²)
 - Indicates percentage of variance in house prices explained by features

Higher R² and lower MSE indicate better model performance.

7. Hyperparameter Tuning:

Standard Linear Regression has no tunable hyperparameters.

To improve performance, **Ridge Regression (L2 Regularization)** can be applied.

- Tuned Parameter: alpha

- Tested Values: [0.1, 1, 10]
- Helps reduce overfitting and improve generalization.

Logistic Regression:

1. Dataset Source:

- **Dataset Name:** Weather Forecast Dataset
- **Source:**
<https://www.kaggle.com/datasets/zeeshier/weather-forecast-dataset>
- **Repository Owner:** Zeeshier

2. Dataset Description:

This dataset contains **2,500 weather observations** and is designed for beginners to practice **classification using machine learning**.

The dataset is used to predict whether **rainfall will occur** based on several atmospheric and environmental conditions.

- **Size:** 2,500 samples (rows) × multiple attributes (columns)
- **Target Variable:**
 - Rain (Binary)
 - 1 → Rainfall occurs
 - 0 → No rainfall
- **Key Features:**

1. Temperature – Current temperature
2. Humidity – Moisture content in air
3. WindSpeed – Speed of wind
4. Pressure – Atmospheric pressure
5. CloudCover – Percentage of cloud cover
6. Visibility – Distance that can be clearly seen
7. DewPoint – Temperature at which condensation occurs

The dataset is **clean, simple, and beginner-friendly**, making it suitable for experimenting with classification algorithms such as Logistic Regression, Decision Trees, and Random Forests.

3. Mathematical Formulation of the Algorithm:

Logistic Regression estimates the probability that an instance belongs to a particular class (Rain or No Rain).

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Where:

$$z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

Decision Rule:

$$\hat{y} = \begin{cases} 1, & \text{if } \sigma(z) \geq 0.5 \\ 0, & \text{if } \sigma(z) < 0.5 \end{cases}$$

4. Algorithm Limitations:

1. **Linear Decision Boundary:** Cannot model complex non-linear relationships.
2. **Sensitive to Feature Scaling:** Features should be standardized.
3. **Performance Drops with Highly Correlated Features.**

5. Methodology / Workflow:

1. Load weather dataset
2. Separate features and target (Rain)
3. Apply feature scaling (StandardScaler)
4. Split dataset into training and testing sets (80:20)
5. Train Logistic Regression model
6. Predict rainfall
7. Evaluate performance

6. Performance Analysis:

- **Metric 1: Accuracy**
 - Measures overall correctness of predictions

- **Metric 2: Confusion Matrix**

- Shows True Positives, False Positives, True Negatives, False Negatives

Higher accuracy and balanced confusion matrix indicate better model performance.

7. Hyperparameter Tuning:

- Parameter Tuned: **C** (Inverse of Regularization Strength)
- Tested Values: [0.01, 0.1, 1, 10, 100]

Impact:

- Low C → Strong regularization (simpler model)
- High C → Weak regularization (complex model)

Code And Output:

```
from google.colab import files
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
uploaded = files.upload()

df = pd.read_csv(list(uploaded.keys())[0])

X = df.drop('SalePrice', axis=1)

y = df['SalePrice']

X = X.select_dtypes(include=[np.number])

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

model = LinearRegression()

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

mae = metrics.mean_absolute_error(y_test, y_pred)

mse = metrics.mean_squared_error(y_test, y_pred)

rmse = np.sqrt(mse)

r2 = metrics.r2_score(y_test, y_pred)

mape = np.mean(np.abs((y_test - y_pred) / y_test))

print("---- Model Performance Metrics ----")

print(f"Mean Absolute Error (MAE) : {mae:.4f}")

print(f"Mean Squared Error (MSE) : {mse:.4f}")

print(f"Root Mean Squared Error (RMSE) : {rmse:.4f}")

print(f"R-squared Score (R2) : {r2:.4f} ({r2*100:.2f}%)")

print(f"Mean Absolute Percentage Error (MAPE) : {mape:.4f} ({mape*100:.2f}%)")
```

```

print("\n---- Regression
Equation Details ----")

print(f"Intercept (b0) :
{model.intercept_:.4f} ")

coeff_df =
pd.DataFrame(model.coef_,
X.columns,
columns=["Coefficient"])

print(coeff_df)

equation = f"SalePrice =
{model.intercept_:.4f}"

for i, col in enumerate(X.columns):
    equation += f" + ({model.coef_[i]:.4f} * {col}) "

```

```

print("\nMathematical
Equation:")

print(equation)

plt.figure(figsize=(8,6))

plt.scatter(y_test, y_pred,
edgecolor='k', alpha=0.7)

plt.plot([y.min(), y.max()],
[y.min(), y.max()])

plt.xlabel("Actual Sale
Price")

plt.ylabel("Predicted Sale
Price")

plt.title(f"Actual vs
Predicted Sale Price
(Accuracy: {r2*100:.2f}%)")

plt.grid(True)

plt.show()

```

---- Model Performance Metrics ----

Mean Absolute Error (MAE): 10126.1733

Mean Squared Error (MSE): 144540866.0004

Root Mean Squared Error (RMSE): 12022.5150

R-squared Score (R2): 0.9770 (97.70%)

Mean Absolute Percentage Error (MAPE): 0.0470 (4.70%)

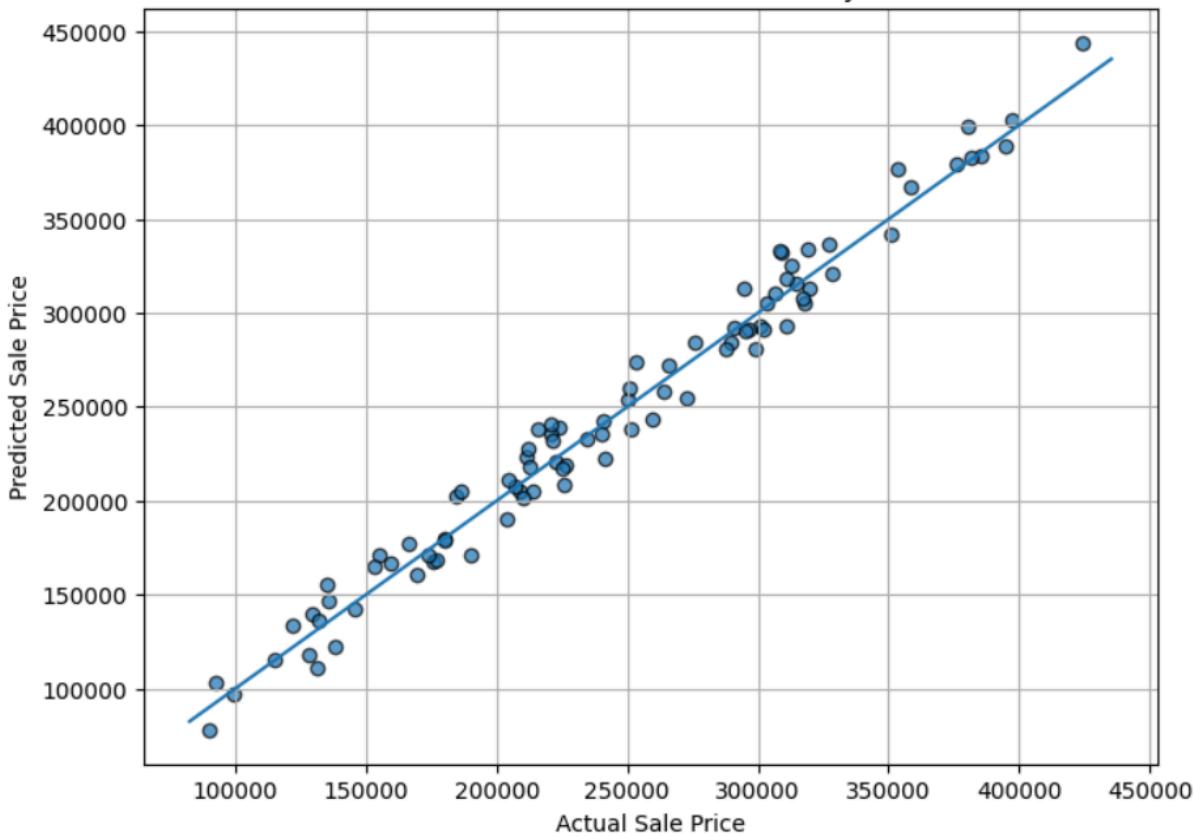
---- Regression Equation Details ----

Intercept (b_0): -1555960.9020

Coefficient

Id	3.695564
OverallQual	12033.078106
GrLivArea	55.247721
GarageCars	19399.549319
TotalBsmtSF	25.489980
YearBuilt	767.411834
FullBath	10746.959166
BedroomAbvGr	4968.399174
LotArea	-0.203670

Actual vs Predicted Sale Price (Accuracy: 97.70%)



```
from google.colab import files
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score,
                           confusion_matrix,
                           precision_score,
                           recall_score, f1_score
uploaded = files.upload()
df = pd.read_csv(list(uploaded.keys())[0])
target = df.columns[-1]
```

```
if df[target].dtype == object:
    df[target] = df[target].astype('category')
    .cat.codes
X = df.drop(target, axis=1)
y = df[target]
X = X.select_dtypes(include=[np.number])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```

model = LogisticRegression(max_iter=10000)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

acc = accuracy_score(y_test, y_pred)

prec = precision_score(y_test, y_pred)

rec = recall_score(y_test, y_pred)

f1 = f1_score(y_test, y_pred)

print("---- Model Performance Metrics ----")

print(f"Accuracy: {acc:.4f} ({acc*100:.2f}%)")

print(f"Precision: {prec:.4f} ({prec*100:.2f}%)")

print(f"Recall: {rec:.4f} ({rec*100:.2f}%)")

print(f"F1 Score: {f1:.4f} ({f1*100:.2f}%)")

cm = confusion_matrix(y_test, y_pred)

tn, fp, fn, tp = cm.ravel()

print(f"TP: {tp}, TN: {tn}, FP: {fp}, FN: {fn}")

plt.figure(figsize=(6,5))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')

plt.xlabel("Predicted Label")

plt.ylabel("True Label")

plt.title("Confusion Matrix")

plt.show()

z = model.decision_function(X_test)

prob = model.predict_proba(X_test)[:, 1]

```

```
plt.figure(figsize=(8,6))

idx = np.argsort(z)
plt.scatter(z, prob,
alpha=0.3)

plt.plot(z[idx], prob[idx])

plt.axhline(0.5)

plt.xlabel("Linear Decision Function (z)")
```

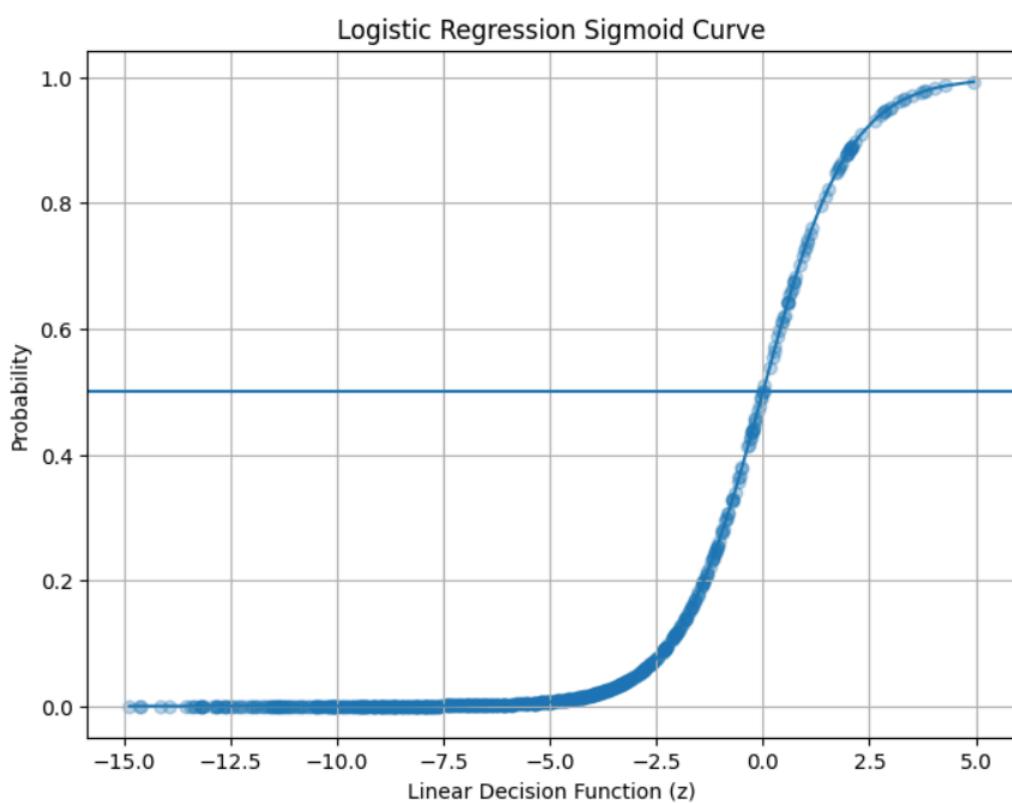
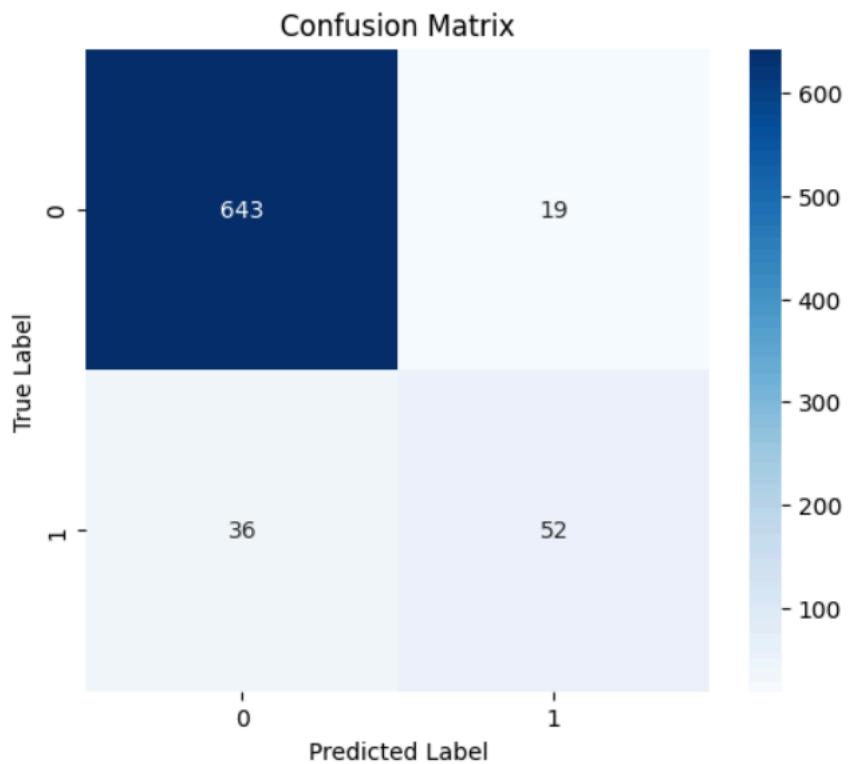
```
plt.ylabel("Probability")

plt.title("Logistic Regression Sigmoid Curve")

plt.grid(True)

plt.show()
```

---- Model Performance Metrics ----
Accuracy: 0.9267 (92.67%)
Precision: 0.7324 (73.24%)
Recall: 0.5909 (59.09%)
F1 Score: 0.6541 (65.41%)
TP: 52, TN: 643, FP: 19, FN: 36



Conclusion:

This experiment demonstrated the practical implementation of Linear Regression for house price prediction and Logistic Regression for rainfall prediction using Kaggle datasets. Both models successfully captured relationships between input features and their respective target variables, showing the effectiveness of supervised learning techniques for solving real-world problems.

The experiment highlighted the importance of data preprocessing, feature selection, feature scaling, and train–test splitting in building reliable machine learning models. Performance evaluation using metrics such as Mean Squared Error, R^2 score, Accuracy, and Confusion Matrix provided insights into model effectiveness. Overall, this experiment strengthened understanding of regression and classification algorithms and their application in predictive analytics.